

An Algorithm for Reversible Logic Circuit Synthesis Based on Tensor Decomposition

Abstract

An algorithm for reversible logic synthesis is proposed. The task is, for given n -bit substitution map $P_n : \{0, 1\}^n \rightarrow \{0, 1\}^n$, to find a sequence of reversible logic gates that implements the map. The gate library adopted in this work consists of multiple-controlled Toffoli gates denoted by C^mX , where m is the number of control bits that ranges from 0 to $n - 1$. Controlled gates with large m (> 2) are then further decomposed into C^0X , C^1X , and C^2X gates. A primary concern in designing the algorithm is to reduce the use of C^2X gate (also known as Toffoli gate) which is known to be universal [29].

The main idea is to view an n -bit substitution map as a rank- $2n$ tensor and to transform it such that the resulting map can be written as a tensor product of a rank- $(2n - 2)$ tensor and the 2×2 identity matrix. Let \mathcal{P}_n be a set of all n -bit substitution maps. What we try to find is a size reduction map $\mathcal{A}_{\text{red}} : \mathcal{P}_n \rightarrow \{P_n : P_n = P_{n-1} \otimes I_2\}$. One can see that the output $P_{n-1} \otimes I_2$ acts nontrivially on $n - 1$ bits only, meaning that the map to be synthesized becomes P_{n-1} . The size reduction process is iteratively applied until it reaches tensor product of only 2×2 matrices.

Time complexity of the algorithm is exponential in n as most previously known algorithms for reversible logic synthesis also are, but it terminates within reasonable time for not too large n which may find practical uses. As stated earlier, our primary concern is the number of Toffoli gates in the output circuit or *quality* for short, not the time complexity of the algorithm. Benchmark results show that the quality of circuits obtained in this work outperforms that of the previously reported ones for almost all hard benchmark functions suggested in [24, 17]. A working code written in Python is publicly available from GitHub [1]. The algorithm is also applied to find reversible circuits for cryptographic substitution boxes which are being used in a certain type of encryption algorithms.

1 Introduction

Beginning from the mid-twentieth century, studies on reversible computing have been motivated by several factors such as power consumption, debugging, routing, performance issues in certain cases, and so on [26]. The circuit-based quantum computing model is also closely related to reversible computing, where every component of the circuit works as a unitary transformation except the measurement [21]. As in classical logic synthesis where abstract circuit behavior is designed by using a specified set of logic gates such as {AND, NOT}, quantum logic synthesis is a process of designing a logic circuit in terms of certain reversible gates. Early studies on quantum algorithms have been carried out in a rather abstract fashion, for example by adopting a unitary oracle specifying only its inputs and outputs, but recent progress in quantum computing and related fields seek lower-level descriptions involving the logic synthesis. For example in quantum cryptanalysis, a number of research groups are trying to measure the complexity of quantum algorithms in terms of elementary quantum gates [13, 15, 33].

Bijection maps of the form $P_n : \{0, 1\}^n \rightarrow \{0, 1\}^n$ often appears as a target behavior to be synthesized in various computational problems, for example in analyzing cryptographic substitution boxes (S-box) [15] or hidden weighted bit functions [6, 5], or permutation network routing [7]. Let us call them permutation maps. Reversible logic synthesis on such maps has been studied intensively [18, 27, 31]. For more relevant works, see, the benchmark pages [24] and [17], and related references therein. It is especially a nontrivial problem to synthesize a reversible circuit for permutation behavior that does not have an apparent structure. Most known algorithms for unstructured permutations with $n > 4$ are heuristic ones with the runtime exponential in n .

One of the well-known methods for synthesizing an n -bit permutation is to find a process of transforming P_n into $\sigma_{n,\text{id}}$, where $\sigma_{n,\text{id}}$ is the n -bit identity permutation. See, Section 2.2 for the reason why finding the process is equivalent to design a reversible circuit. A straightforward way to achieve the goal is to rewrite the permutation as a product of at most $2^n - 1$ transpositions, and to synthesize each transposition in terms of specified gates [21]. This naive approach is easy to implement, but the resulting circuit involves a large number of Toffoli gates. Researchers have introduced various algorithms to improve the method, and what we have noticed is that instead of finding a direct map for identity, one may try a map that reduces the effective size of the permutation such that $P_n \mapsto P_{n-1}$ and iteratively apply it. Somewhat related notion has been examined in synthesizing unitary matrices [19, 25], but apart from the use of the notion *block* (see, Section 3.1), the design criteria are quite different. In a nutshell, their method is to divide a large circuit into several smaller circuits and each one is recursively divided into even smaller ones, involving gate operations to every bit all the way through to the end. On the contrary, each time the size of the permutation gets smaller in our method, one bit becomes completely irrelevant from the circuit.

This work is summarized as follows:

- An algorithm for reversible logic circuit synthesis is proposed. The algorithm is designed so that the output circuit involves as small number of Toffoli gates as possible. Comparisons with the best results previously reported in [24, 17] are summarized in Table 1.
- The algorithm is applied to AES [23], Skipjack [22], KHAZAD [4], and DES [20] S-boxes. The point of applying the algorithm to the S-boxes is that except AES, all other S-boxes do not have apparent structures where no known polynomial-time algorithm can be applied.

An easy-to-use implementation of the algorithm written in Python is also available from GitHub [1].

The paper is organized as follows. Section 2 briefly covers the basics of reversible logic circuit synthesis and relevant works. Sections 3 and 4 are devoted to bringing out the design criteria of the algorithm. Benchmark results and applications to cryptographic S-boxes are presented in Section 5.

Table 1: Benchmark results for hard benchmark functions from [24, 17]. The first three subcolumns read the number of input, output, and garbage bits, respectively. QC stands for Quantum Cost which is one of the widely accepted cost metrics [2], and #TOF is the number of Toffoli gates.

Functions	Best reported					This work				
	#in	#out	#grb	QC	#TOF	#in	#out	#grb	QC	#TOF
URF1	9	9	0	17002	2225	9	9	0	13318	2029
URF2	8	8	0	7083	892	8	8	0	5510	803
URF3	10	10	0	37517	4997	10	10	0	33541	4898
URF4	11	11	0	160020	32004	11	11	0	81748	12516
URF5	9	9	0	14549	1964	9	9	0	11902	1527
<i>nthPrime7</i>	7	7	0	2284	334	7	7	0	1984	292
<i>nthPrime8</i>	8	8	0	6339	930	8	8	0	5333	726
<i>nthPrime9</i>	10	9	1	17975	2392	9	9	0	13918	2073
<i>nthPrime10</i>	11	10	1	40299	5302	10	10	0	30421	4357
<i>nthPrime11</i>	12	11	1	95431	12579	11	11	0	69702	10366




Figure 1: Truth tables and circuit symbols for CNOT (left) and Toffoli (right) gates.

2 Backgrounds

Logic circuit synthesis involves a functionally complete set of logic gates such as {AND, NOT}, which we would call a universal gate set. In this section, we first briefly introduce widely adopted *reversible* gate sets, which are called multiple-controlled Toffoli (MCT) library and NOT, CNOT, Toffoli (NCT) library [24].

While reversible computing deals with any reversible operations, here we focus only on n -bit input and n -bit output logic behaviors. Therefore we will narrow our attention to the permutation circuit synthesis problem defined in Section 2.2.

Throughout the paper, the ordinal numbers usually begin with zeroth, except for the position of a bit. In referring to the position of a bit, it begins with the ‘first’ for the most significant bit and ends with the ‘ n -th’ for the least significant bit, not $(n - 1)$ -th.

2.1 Gate Library Controlled-NOT, also known as CNOT, takes two input bits, one being control and the other being a target. Similarly, a Toffoli gate takes three input bits, two being controls and the other being a target. See, Fig. 1 for CNOT and Toffoli gates which are clearly reversible. NOT gate is already reversible, and NCT library is known to be a universal set, meaning that any reversible logic can be implemented by using these gates [29].

Naturally, CNOT can be generalized to take $m + 1$ input bits with m controls and one target. It




Figure 2: Decomposition of C^4X into five ($= 2 \cdot 4 - 3$) Toffoli gates. Here a, b, c, d are input control bits and w_0, w_1 are zeroed work bits.

is sometimes called an MCT gate [24]. Let C^mX denote such gates where m is a non-negative integer. Since NCT gates are already included in MCT library (as with $m = 0, 1, 2$), MCT library is also a universal set.

A C^mX gate with $m > 2$ can be decomposed into a few NCT gates with various tradeoffs. One may think of it as a conversion formula from an MCT gate to NCT gates. For example, C^mX can be constructed by using $2m - 3$ Toffoli gates with $m - 2$ zeroed work bits, or by using $8m - 24$ Toffoli gates with one arbitrary work bit [3]. Fig. 2 illustrates the former formula.

The number of Toffoli gates presented in a circuit will be regarded as the *quality* of it. Note that the quality gets better as the number of Toffoli gates decreases. An algorithm we design primarily uses MCT library, but then to measure the quality of the output, we need a certain conversion formula for all C^mX gates with $m > 2$. For this purpose, we will use a simple formula described in Fig. 2.

2.2 Permutation Circuit Synthesis Having an appropriate basis, an n -bit permutation written in one-line notation $(r_0, r_1, \dots, r_{2^n-1})$ can be seen as a matrix in $\{0, 1\}^{2^n \times 2^n}$. For example for $n = 3$ with the standard basis $\mathbf{b}_0 = (1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0)^\top$, $\mathbf{b}_1 = (0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0)^\top, \dots$, a permutation $P_3 = (7, 2, 0, 1, 5, 3, 6, 4)$ can be written by a truth table and by a matrix

$$\begin{array}{c|c}
 \text{in} & \text{out} \\
 \hline
 000 & 111 \\
 001 & 010 \\
 010 & 000 \\
 011 & 001 \\
 100 & 101 \\
 101 & 011 \\
 110 & 110 \\
 111 & 100
 \end{array}
 , \quad
 P_3 = \begin{pmatrix}
 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\
 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\
 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\
 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0
 \end{pmatrix}, \quad (2.1)$$

where we have used the same symbol P_n interchangeably to denote the one-line notation of an n -bit permutation and the corresponding matrix. Having it in mind, we define permutation circuit synthesis as follows:

DEFINITION 2.1. For an n -bit permutation P_n , a finite universal gate set G , and a cost metric on a set, permutation circuit synthesis is to find a finite ordered set $R = \{g_i \mid g_i \in G\}$ such that $P_n \cdot (g_1 \cdot g_2 \cdot \dots \cdot g_{|R|}) = I_{2^n}$, minimizing the cost on R .

Denoting NOT gate acting on i -th bit by X_i , CNOT gate controlled by i -th bit acting on j -th

bit by CX_{ij} , and Toffoli gate controlled by i - and j -th bits acting on k -th bit by CX_{ijk} , one can verify that $P_3 \cdot CX_{13} \cdot CX_{312} \cdot X_2 \cdot CX_{23} \cdot CX_{231} = I_{2^3}$ (see, Eq (3.3) for gate actions), and thus $CX_{231} \cdot CX_{23} \cdot X_2 \cdot CX_{312} \cdot CX_{13} = P_3$ since NCT gates are involutory ($A = A^{-1}, A \in \text{NCT}$). At this point, a circuit for P_3 is synthesized but it can be questioned if the circuit is optimal under a certain cost metric. In fact, the optimality of 3- and 4-bit permutations under certain cost metrics has been proved by using exhaustive search [16] or meet-in-the-middle technique [9], but for larger n no known result has claimed it.

When a permutation is structured it is often possible to find an efficient circuit by exploiting its structure. Here by ‘structured’ we mean r_i is efficiently computable such that $r_i = f(i)$ with some function f . If such function is known, one may design a circuit that computes f reversibly. A good example would be an eight-bit S-box used in the symmetric block cipher AES algorithm that can be calculated by arithmetic operations [23, 10].

On the other hand, when a permutation shows no apparent structure there is no known efficient method to synthesize it. The remaining options are to apply heuristic algorithms [26] or reversible lookup table methods [28]. A reversible lookup table is to implement each r_i for corresponding i one-by-one, reversibly. Although straightforward and possibly not less efficient than heuristic algorithms asymptotically, the reversible lookup table method is not preferred over heuristic algorithms for the problems at hand to our consideration.

A large number of algorithms for reversible logic circuit synthesis have been proposed. A good review can be found in [26]. Due to the heuristic nature of the algorithms, benchmark pages have been established for performance comparisons [24, 17].

3 Basic Idea

The basic strategy this work has taken is to reduce the effective size of a matrix at each step by looking for a size reduction process, leaving one bit completely excluded from subsequent steps. The size reduction idea can then naturally be resulted in an algorithm for permutation circuit synthesis under a certain assumption. The assumption will be lifted in Section 4 leading to an algorithm that works for any permutation.

3.1 Matrix Size Reduction Reversible gates acting on n bits can be viewed as rank- $2n$ tensors [30]. Among them, there exist certain gates of which decomposition as a tensor product of smaller tensors can be found. For example, Walsh-Hadamard transformation acting on two bits can be represented as a 4×4 matrix (or equivalently rank-4 tensor), which can also be decomposed into two 2×2 matrices such as

$$\begin{aligned} H^{\otimes 2} &= \frac{1}{2} \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & -1 \\ 1 & 1 & -1 & -1 \\ 1 & -1 & -1 & 1 \end{pmatrix} \\ &= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \otimes \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}. \end{aligned} \tag{3.2}$$

However, CNOT cannot be written in terms of a simple tensor product of two smaller tensors.

Now consider a permutation P_n that can be written as $P_{n-1} \otimes I_2$. It means the gate P_n acts nontrivially on $n - 1$ bits only, effectively leaving one bit irrelevant from the operation. It is clear that if one is able to find a procedure for transforming an arbitrary rank- $2n$ tensor into one that can be

decomposed into one rank- $(2n - 2)$ tensor and one 2×2 identity matrix, the circuit can be synthesized by iterating the procedure at most $n - 1$ times.

3.2 Definitions and Conventions In n -bit permutation circuit synthesis, a gate set we use consists of C^mX gates, where the number of control bits ranges from 0 to $n - 1$.

Dealing with a permutation is probably best comprehensible in matrix representation as in Eq. (2.1), with an obvious downside that it is inconvenient to write down. One-line notation is thus frequently adopted throughout the paper. For example, the action of X_1 , CX_{21} , and CX_{312} on $(7, 2, 0, 1, 5, 3, 6, 4)$ reads

$$\begin{aligned} & \xrightarrow{X_1} (5, 3, 6, 4, 7, 2, 0, 1), \\ (7, 2, 0, 1, 5, 3, 6, 4) & \xrightarrow{CX_{21}} (7, 2, 6, 4, 5, 3, 0, 1), \\ & \xrightarrow{CX_{312}} (7, 2, 0, 1, 5, 4, 6, 3). \end{aligned} \tag{3.3}$$

When a permutation is written by $P_n = (r_0, r_1, \dots, r_{2^n-1})$, the subscripts i and integers r_i are understood as *column numbers* and *row numbers*, respectively, in which nonzero values reside in the matrix representation. For example, $r_0 = 7$ means 1 is located at the 0th column and the 7th row in the matrix as in Eq. (2.1). Using these notions, a permutation P_n can be viewed as a function of a column number,

$$\begin{aligned} P_n : \{0, 1, \dots, 2^n - 1\} & \rightarrow \{0, 1, \dots, 2^n - 1\}, \\ i & \mapsto r_i. \end{aligned}$$

Column numbers will frequently be read as n -bit binary strings. Binary strings will be denoted by vector notation when necessary. In writing an integer x as an n -bit binary string $\mathbf{x} \in \{0, 1\}^n$, we let $x_i \in \{0, 1\}$ denotes the i -th bit of \mathbf{x} for $1 \leq i \leq n$.

One way to understand the action of a logic gate is to think of it as an operator that exchanges column numbers. When a gate is applied, it first reads column numbers as n -bit binary strings. Among the strings (columns), some must meet the condition for the activation of the gate. The row numbers reside in these columns are swapped appropriately. For example in Eq. (3.3), each column number is read as 000, 001, \dots , 111 which are occupied by row numbers 7, 2, \dots , 4, respectively. The gate CX_{132} is activated when the value of the first and the third bits are both 1, which correspond to the 5th (101) and the 7th (111) columns. The row numbers 3 and 4 which preoccupied these positions are then swapped by the gate.

In addition to X_i , CX_{ij} , and CX_{ijk} gates defined in Section 2.2, we introduce a way to specify control and target bits of general C^mX gates with nonnegative integer m . Let \mathcal{I} be a subset of $\{1, \dots, n\}$. In denoting a controlled gate with an arbitrary number of control bits, an index set \mathcal{I} will be used such that $CX_{\mathcal{I};k}$ has $|\mathcal{I}|$ control bits specified by elements in \mathcal{I} and targets k -th bit.

The following definition for *block* is the central notion in this work, and it is recommended to refer to the example in Eq. (3.4) before comprehending the formal definition.

DEFINITION 3.1. For an n -bit permutation (r_0, \dots, r_{2^n-1}) , a pair of numbers r_{2i} and r_{2i+1} is defined as an even (odd) block if $r_{2i+1} - r_{2i} = 1$ (-1), where $i \in \{0, 1, \dots, 2^{n-1} - 1\}$ is called a block-wise position.

It is unlikely that the number of blocks found in an unstructured permutation is large. It is then our task to find a transformation to maximize the number. For example, assume there is a map applied to

a permutation matrix in Eq. (2.1) as follows:

$$\begin{pmatrix} 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \mapsto \begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}. \quad (3.4)$$

The 2×2 substructures in gray color are called blocks. Even (odd) blocks are 2×2 identity (off-diagonal) structures. In one-line notation of the resulting matrix $(2, 3, 7, 6, 4, 5, 1, 0)$, pairwise row numbers 2, 3 and 4, 5 are even blocks and 7, 6 and 1, 0 are odd blocks.¹

These pairs of row numbers play an important role hereafter, but there is a chance the notions confuse readers. To avoid confusion, let us describe a way to construct an even block ‘2, 3’ beginning from $(7, 2, 0, 1, 5, 3, 6, 4)$. Using the r_i notation, initially we have the pair $r_1 = 2$ and $r_5 = 3$. First, we want to put 2 in the 0th column. Applying $X_1 C X_{13} X_1$ leads to $(2, 7, 1, 0, 5, 3, 6, 4)$. This procedure can be interpreted as moving row number 2 from column position 1 to 0; $r_1 \mapsto r_0$. Similarly, applying $C X_{31}$ then results in $(2, 3, 1, 4, 5, 7, 6, 0)$; $r_5 \mapsto r_1$. As already explained, applying gates can be interpreted as exchanging column numbers for fixed row numbers. All the algorithms and subroutines introduced below have similar descriptions that, first targeting a certain pair of row numbers, and then changing their column positions. At this point, let us formally define such pairs.

DEFINITION 3.2. *A pair of row numbers $2j, 2j + 1$ for $0 \leq j \leq 2^{n-1} - 1$ is called a relevant row pair, denoted by $\langle 2j, 2j + 1 \rangle$.*

Row numbers comprising a relevant row pair are called relevant row numbers. Constructing a block is thus putting relevant row numbers appropriately side-by-side that are originally located apart.

In addition, we further define the following:

DEFINITION 3.3. *For a relevant row pair $\langle r_i, r_j \rangle$, the row numbers r_i and r_j are said to be occupied at*

- *normal positions if $r_i \equiv i$ and $r_j \equiv j \pmod{2}$.*
- *inverted positions if $r_i \not\equiv i$ and $r_j \not\equiv j \pmod{2}$.*
- *interrupting positions otherwise.*

To better understand the definitions, consider a permutation $(\boxed{7}, \boxed{2}, \boxed{0}, \boxed{1}, \boxed{5}, \boxed{3}, \boxed{6}, \boxed{4})$ and examine the relevant row pairs. Relevant row numbers $r_2 = 0$ and $r_3 = 1$ are at normal positions as $2 \equiv_2 0$ and $3 \equiv_2 1$. Another relevant numbers $r_7 = 4$ and $r_4 = 5$ are at inverted positions as $7 \not\equiv_2 4$ and $4 \not\equiv_2 5$. Other numbers are at interrupting positions. In simple terms, 0 and 1 are at normal positions as the small one is in the gray box, 4 and 5 are at inverted positions as the small one is in the white box, and other numbers are at interrupting positions as the numbers are in the same colored boxes.

It may seem plausible that the relevant numbers in normal (inverted) positions in the beginning likely end up in even (odd) blocks in the series of transformations for maximizing the number of blocks. Indeed if we are careful enough, all the row numbers in the normal (inverted) positions in the beginning form even (odd) blocks at the end. A formal description for the observation is as follows:

¹Note that the resulting matrix in Eq. (3.4) is not a tensor product of two smaller tensors, yet. Further application of $C X_{23}$ achieves four even blocks thereby completing the procedure.

REMARK 3.1. When C^mX gate is applied, the number of normal, inverted, and interrupting positions is conserved unless the gate is targeting the n -th bit.²

For example, a permutation (7, 2, 0, 1, 5, 3, 6, 4) has four row numbers at interrupting positions; 2, 3 and 7, 6. The action of CX_{32} that does not *target* the 3rd bit leads to the permutation (7, 1, 0, 2, 5, 4, 6, 3), which still has four interrupting positions. On the other hand, if CX_{13} is applied, the resulting permutation (7, 2, 0, 1, 3, 5, 4, 6) will no longer involve any interrupting position, but instead, the number of normal and inverted positions will be increased by two each. Details will not be covered, but we would point out that handling the interrupting positions is typically more expensive than dealing with normal or inverted positions. Therefore in the next section, a preprocess that removes interrupting positions before getting into the main process will be introduced. A way to deal with the inverted positions will also be introduced in the next section. For now, let us restrict our attention to permutations that only have normal positions. It will soon turn out that reversible gates targeting the last bit are not required at all in this section.

A natural way to design an algorithm for the size reduction could be to build blocks one-by-one, without losing already built ones. Remind that in one-line notation for a given permutation $(\boxed{r_0, r_1}, \boxed{r_2, r_3}, \boxed{r_4, r_5}, \boxed{r_6, r_7}, \dots)$, a block is a pair of row numbers r_i that differ by 1 (with the odd one being larger) and both reside in one of the designated column positions (boxes). Applying reversible gate swaps positions of at least two row numbers as in Eq. (3.3). What we want to do is to construct a new block while maintaining already constructed ones, which very much resembles solving Rubik's Cube. Each rotation in Rubik's Cube corresponds to the application of a logic gate. The difference is that while the number of rotations is to be minimized typically in Rubik's Cube, we try to minimize the number of Toffoli gates.

3.3 Glossary All necessary definitions and notations have been introduced, which we summarize below.

X_i, CX_{ij}, CX_{ijk}	NCT gates (Section 2.2)
$CX_{\mathcal{I}:k}$	MCT gates (Section 2.1)
R	An ordered set of gates
P_n	An n -bit substitution map
$(r_0, r_1, \dots, r_{2^n-1})$	One-line notation of P_n
Row number r_i	Below Eq. (3.3)
Column number i (of r_i)	Below Eq. (3.3)
Block	Definition 3.1
Block-wise position	Definition 3.1
Relevant row pair	Definition 3.2
Relevant row numbers	Below Definition 3.2
Normal position	Definition 3.3
Inverted position	Definition 3.3
Interrupting position	Definition 3.3
\mathbf{x}	Binary string of an integer x
x_i	The i -th bit of \mathbf{x} for $1 \leq i \leq n$

3.4 Size Reduction Algorithm In this section we restrict our attention to a procedure $\mathcal{A}'_{\text{red}} : \mathcal{P}'_n \rightarrow \{P_n : P_n = P_{n-1} \otimes I_2\} \times \mathcal{R}$, where \mathcal{P}'_n is a set of all n -bit permutations of which row numbers are only

²More specifically, each row number remains still in its respective position unless the specified gates are involved, but we only need the fact that the number of such positions is conserved.

in normal positions and \mathcal{R} is a set of all ordered set of MCT gates. The algorithm $\mathcal{A}'_{\text{red}}$ is designed obeying three rules.

- A block is constructed one-by-one.
- The constructed block is allocated to the left in one-line notation.
- The number of left-allocated blocks should not decrease upon the action of any gate.

The meaning of the construction and the allocation is given in an example. Assume a three-bit permutation $P_3 = (0, 1, 6, 3, 2, 5, 4, 7)$ is at hand. The first block is already there, and we target the next block $(4, 5)$. Applying CX_{312} results in $P_3 \cdot CX_{312} = (0, 1, 6, 3, 2, 7, \boxed{4, 5})$. One can see that a new block is *constructed* in the 6th and the 7th columns at the cost of one Toffoli gate. The constructed block is then *allocated* in the 2nd and the 3rd columns by applying CX_{21} , i.e., $P_3 \cdot CX_{312} \cdot CX_{21} = (0, 1, \boxed{4, 5}, 2, 7, 6, 3)$. In this example, only the construction costs a Toffoli gate whereas the allocation does not, but in general allocation also costs Toffoli gates. In spite of the seemingly unnecessary cost for the allocation, we conclude that the left-allocation is more beneficial than leaving a block where it is constructed. Detailed reasonings for the left-allocation will not be covered³.

In describing algorithms and subroutines, operations that update the permutation or the gate sequence will be frequently involved. For a permutation P_n , an ordered set of gates $R = (g_1, \dots, g_{|R|})$, and an another ordered set of gates $S = (h_1, \dots, h_{|S|})$, define a symbol \cdot such that

$$(P_n, R) \cdot S = (P_n \cdot h_1 \cdot h_2 \cdots h_{|S|}, (R; S)),$$

$$(R; S) = (g_1, \dots, g_{|R|}, h_1, \dots, h_{|S|}).$$

High-level description of the size reduction algorithm is as follows:

Algorithm 1 $\mathcal{A}'_{\text{red}}$

Input P_n

$\triangleright P_n \in \mathcal{P}'_n$

Output (P, R)

$\triangleright P = P_{n-1} \otimes I_2$

Procedure

- 1: $R \leftarrow ()$ \triangleright An empty ordered set
 - 2: $P \leftarrow P_n$
 - 3: **for** i from 0 to $2^{n-1} - 1$ **do**
 - 4: $\langle a, b \rangle \leftarrow \text{PICK}(P, i)$ \triangleright Pick a relevant row pair
 - 5: $S_C \leftarrow \text{CONS}(P, i, \langle a, b \rangle)$
 - 6: $(P, R) \leftarrow (P, R) \cdot S_C$
 - 7: $S_A \leftarrow \text{ALLOC}(P, i, a)$
 - 8: $(P, R) \leftarrow (P, R) \cdot S_A$
 - 9: **return** (P, R)
-

It will be helpful to have some intuition behind the details before describing the subroutines PICK,

³Roughly explained, if the blocks are located randomly across the possible positions, it will get more and more difficult to construct a new block while maintaining the already constructed ones since cheaper logic gates tend to stir many positions. One way to mitigate the difficulty is to make blocks share the same bit value in the binary string of the column numbers so that certain controlled-gates can leave such blocks unaffected by using that bit as a control.

and a few CX and X gates, where $s(i)$ is a function such that $s(i) \leq HW(i)$, where $HW(i)$ is the hamming weight of i .

We first describe PICK.

Subroutine 1 PICK

Input P_n, i

Output $\langle a, b \rangle$

Procedure

```

1:  $m \leftarrow \text{FINDM}(i)$ 
2:  $k \leftarrow 2^n - 2^{n-m+1}$ 
3: for  $j$  from  $k$  to  $2^n - 2$  do
4:    $a \leftarrow r_j$ 
5:    $b \leftarrow a \oplus \mathbf{1}$   $\triangleright \mathbf{1} = 00 \dots 01$ 
6:   for  $t$  from  $j + 1$  to  $2^n - 1$  do
7:      $c \leftarrow r_t$ 
8:     if  $b = c$  then
9:       return  $\langle a, b \rangle$ 

```

Termination of the subroutine will be shown to be guaranteed in the next subsection. In line 1, FINDM computes m from i by letting $l = i$ in Eq. (3.5).

Now we describe CONS.

Subroutine 2 CONS

Input $P_n, i, \langle a, b \rangle$ **Output** S_C **Procedure**

```
1:  $S_C \leftarrow ( )$  ▷ An empty ordered set
2:  $m \leftarrow \text{FINDM}(i)$ 
3:  $(\alpha, \beta) \leftarrow \text{COL}(P_n, a, b)$  ▷ Find column numbers
4:  $\gamma \leftarrow \alpha \oplus \beta$ 
5:  $\delta \leftarrow 0$ 
6: for  $j$  from 1 to  $n$  do
7:   if  $\gamma_j = 1$  then
8:      $\delta \leftarrow j$ 
9:     break
10: if  $\delta = n$  then ▷ Already a block
11:   return  $S_C$ 
12: if  $i_\delta = 1$  then
13:    $S_C \leftarrow S_C; X_\delta$  ▷ Append  $X_\delta$  to  $S_C$ 
14: for  $j$  from  $\delta + 1$  to  $n - 1$  do
15:   if  $\gamma_j = 1$  then
16:      $S_C \leftarrow S_C; CX_{\delta j}$ 
17: if  $i_\delta = 1$  then
18:    $S_C \leftarrow S_C; X_\delta$ 
19:  $\mathcal{I} \leftarrow \{n\}$ 
20: for  $j$  from 1 to  $m - 1$  do
21:    $\mathcal{I} \leftarrow \mathcal{I} \cup \{j\}$ 
22:  $S_C \leftarrow S_C; CX_{\mathcal{I}, \delta}$ 
23: return  $S_C$ 
```

In line 3, COL outputs column numbers α, β corresponding to the relevant row numbers a, b . This procedure is necessary since what CONS essentially does is to swap columns as mentioned earlier.

ALLOC is similarly described.

Subroutine 3 ALLOC

Input P_n, i, a **Output** S_A **Procedure**

```

1:  $S_A \leftarrow ( )$ 
2:  $m \leftarrow \text{FINDM}(i)$ 
3:  $\alpha \leftarrow \text{COL}(P_n, a)$ 
4:  $\gamma \leftarrow \alpha \oplus i$ 
5:  $\delta \leftarrow 0$ 
6: for  $j$  from 1 to  $n$  do
7:   if  $\gamma_j = 1$  then
8:      $\delta \leftarrow j$ 
9:     break
10: if  $\delta = 0$  then
11:   return  $S_A$  ▷ Already allocated
12:  $\mathcal{I} \leftarrow \{ \}$ 
13: for  $j$  from  $\delta + 1$  to  $n - 1$  do
14:   if  $\gamma_j = 1$  then
15:      $S_A \leftarrow S_A; CX_{\delta j}$ 
16:   if  $i_j = 1$  then
17:      $\mathcal{I} \leftarrow \mathcal{I} \cup \{j\}$ 
18:  $S_A \leftarrow S_A; CX_{\mathcal{I}:\delta}$ 
19: return  $S_A$ 

```

3.5 Quality and Complexity It has been sketched in Section 3.4 how the size reduction algorithm works. Here we evaluate the time complexity and the quality of the output by introducing two lemmas, one for the construction and the other for the allocation.

PROPOSITION 3.1. *Let l be the number of left-allocated blocks, and m be the smallest positive integer satisfying Eq. (3.5). Define a function $h_n : \mathbb{N} \rightarrow \mathbb{Z}$, $h_n(x) = 2^n - 2^{n-(x-1)}$. There exists at least one relevant row pair r_α, r_β such that*

$$h_n(m) < \min(\alpha, \beta), \quad (3.6)$$

where α, β are column numbers of r_α, r_β , respectively.

Proof. It is nothing more than Pigeonhole principle. It is trivial for $m = 1$. For $m > 1$, there exist $2^n - 2l$ row numbers that do not form left-allocated blocks, yet. Let us call them the remaining row numbers (row pairs). Assume there is no remaining row pair satisfying Eq. (3.6). Then at least one relevant row number in every remaining pair has to sit in column numbers no greater than $h_n(m)$. However, since $2l$ column numbers are already occupied by left-allocated blocks, there only exist $h_n(m) - 2l$ columns available for the row numbers to sit in. If we subtract $h_n(m) - 2l$ from half the number of remaining row numbers $(2^n - 2l)/2$,

$$\begin{aligned} \frac{(2^n - 2l)}{2} - (h_n(m) - 2l) &= l - \frac{h_n(m) - 1}{2} \\ &> 0, \end{aligned}$$




Figure 4: L_δ and R_δ for $\delta = 1, 2, 3$.

where the inequality follows from the fact that $h_n(m-1) < 2l \leq h_n(m)$ by the definition of m . Therefore the assumption cannot be true and there must exist at least one pair that satisfies Eq. (3.6).

LEMMA 3.1. *Let i be an iterator used in Algorithm 1, l be the number of left-allocated blocks, and m be the smallest positive integer satisfying Eq. (3.5). For given permutation, a new block can be constructed by using at most $n - m - 1$ CX , two X , and one C^mX gates, without breaking left-allocated blocks.*

Proof. Let $\mathcal{C} = \{0, 1, \dots, 2^{n-1} - 1\}$ be a set of block-wise positions of an n -bit permutation, i be the binary string of $i \in \mathcal{C}$, and r_α, r_β be relevant row numbers satisfying Eq. (3.6). Viewing their column numbers as binary strings α and β , the first $m - 1$ bits of them are 1 (none if $m = 1$),

$$\alpha_1 = \beta_1 = \alpha_2 = \beta_2 = \dots = \alpha_{m-1} = \beta_{m-1} = 1. \quad (3.7)$$

Let $\gamma = \alpha \oplus \beta$ and δ be the smallest integer such that $\gamma_\delta = 1$. It can be seen that $\delta > m - 1$. Assume $\delta \neq n$, otherwise, the pair is already a block. Let L_δ, R_δ be disjoint subsets of \mathcal{C} such that $\mathcal{C} = L_\delta \cup R_\delta$ and the blocks residing in L_δ and R_δ are preserved under the actions of $CX_{\delta x}$ and $X_\delta CX_{\delta x} X_\delta$ gates for $x \in \{1, \dots, \delta - 1, \delta + 1, \dots, n\}$, respectively. A few L_δ and R_δ are illustrated in Fig. 4.

At i -th iteration, note that i (as a block-wise number) is either included in L_δ or R_δ . If $i \in L_\delta$, an action of $CX_{\delta x}$, $\delta < x < n$ gate does not break left-allocated blocks since the blocks in L_δ are not affected at all and the blocks in R_δ change their respective positions within $< i$. If on the other hand $i \in R_\delta$, similarly an action of $X_\delta CX_{\delta x} X_\delta$, $\delta < x < n$ gate does not break left-allocated blocks since the blocks in R_δ are conserved and the blocks in L_δ change their respective positions within $< i$. Because i_δ tells if it is included in L_δ or R_δ , we are able to apply the aforementioned gates without breaking left-allocated blocks.

First $\delta - 1$ bits of γ are zeros due to the definition of δ . Since $CX_{\delta x}$ or $X_\delta CX_{\delta x} X_\delta$, $\delta < x < n$ gate can be applied without breaking the left-allocated blocks, it is always possible to achieve $\gamma_\delta = \gamma_n = 1$ and $\gamma_k = 0$ for all $k \notin \{\delta, n\}$, without involving any Toffoli gate. Once it is achieved, by applying $CX_{\mathcal{I}:\delta}$, $\mathcal{I} = \{1, 2, \dots, m - 1, n\}$, we have $\gamma = 00 \dots 01$, which is by definition a block. Blocks located left to i are unaffected by $CX_{\mathcal{I}:\delta}$ gate since the binary strings of their positions have at least one zero among the first $m - 1$ bits.

LEMMA 3.2. *Let i be an iterator used in Algorithm 1, l be the number of left-allocated blocks, and m be the smallest positive integer satisfying Eq. (3.5). For given permutation with a block constructed by CONS, the block can be allocated to the i -th block-wise position by using at most $(n - m)$ CX and one $C^{HW(i')}X$ gates, without breaking left-allocated blocks, where $i' = i - h_n(m - 1)/2$ for $m > 1$ and $i' = i$ otherwise.*

Proof. Let \mathbf{j} be a binary string of the (block-wise) position j of the block constructed by CONS. Assume $j \neq i$, otherwise, the block is already left-allocated. Let $\gamma = \mathbf{i} \oplus \mathbf{j}$ and δ be the smallest integer such that $\gamma_\delta = 1$. Due to the property of the column numbers we begin with in CONS, i.e., Eq. (3.6), the first $m - 2$ bits of \mathbf{i} and \mathbf{j} are 1 (none if $m \leq 2$),

$$i_1 = j_1 = i_2 = j_2 = \cdots = i_{m-2} = j_{m-2} = 1, \quad (3.8)$$

and thus $\delta > m - 2$. Since $i < j$ and $\gamma_1 = \gamma_2 = \cdots = \gamma_{\delta-1} = 0$, it is guaranteed that $i_\delta = 0$ and $j_\delta = 1$. Similar to the technique used in Lemma 3.1, \mathbf{j} can be transformed such that $\gamma_\delta = 1$ and $\gamma_k = 0$ for all $k \notin \{\delta\}$ by using $CX_{\delta x}$ gates for $\delta < x < n$. Note that unlike in Lemma 3.1 where α and β both can be transformed by gates, here only \mathbf{j} is allowed to change. Once it is done, by applying $CX_{\mathcal{I}, \delta}$, $\mathcal{I} = \{x \mid x > \delta, i_x = 1\}$, the left-allocation is completed. Note that blocks residing in L_σ , $\sigma \in \{x \mid x < \delta, i_x = 1\}$ are conserved by an action of a gate whose target is δ -th bit, up to changes of block-wise positions within. In addition, by the definition of L_σ , i cannot be included in any of L_σ . Therefore, $CX_{\mathcal{I}, \delta}$ can be applied without breaking the left-allocated blocks.

The quality of the output circuit naturally follows from two lemmas. Be reminded that we have adopted the conversion formula $C^m X : C^2 X = 1 : 2m - 3$. By counting the worst-case number of Toffoli gates in CONS and ALLOC for all iterations, it can be summarized as follows:

THEOREM 3.1. *The number of Toffoli gates in the output R of Algorithm 1 is upper bounded by $\mathcal{N}_c(n) + \mathcal{N}_a(n)$ for $n \geq 3$, where*

$$\begin{aligned} \mathcal{N}_c(n) &= \sum_{i=2}^{n-1} (2i - 3) \cdot 2^{n-i}, \\ \mathcal{N}_a(n) &= \sum_{j=2}^{n-2} \sum_{i=2}^{n-j} (2i - 3) \cdot \binom{n-j}{i}, \end{aligned} \quad (3.9)$$

where $\mathcal{N}_a(3) = 0$.

The time complexity of Algorithm 1 can also be estimated by the lemmas. Assuming all gates exchange 2^n column numbers (overestimation for $C^m X$ for $m > 0$), the time complexity of the algorithm simply reads the total number of gates times 2^n . The maximum number of gates applied in each CONS and ALLOC is $O(n)$, and the number of CONS and ALLOC called in Algorithm 1 is 2^{n-1} each, and thus the asymptotic time complexity reads $O(n2^{2n})$.

4 Algorithm

Key ideas have been introduced in Section 3. The only downside of Algorithm 1 is its inability to decompose arbitrary permutations. This section is therefore mostly devoted to lifting the assumption that an input permutation consists only of row numbers in normal positions. Lifting the assumption costs an extra number of Toffoli gates which is bounded by

$$5 \cdot 2^{n-4} + 2n - 5 + \sum_{i=2}^{n-3} (2i - 3) \cdot \binom{n-3}{i}. \quad (4.10)$$

Lifting the assumption takes several independent steps, which we shall separately deal with in each subsection. Details on various formulas appearing in this section are more or less the same as in the previous section, and thus will mostly be dropped.

4.1 Heuristic Mixing Throughout Section 4 we will frequently refer to the ratio of the number of normal, inverted, and interrupting positions as the ratio $x : y : z$ such that $x + y + z = 1$. For example, all permutations handled in Section 3.4 have the ratio $1 : 0 : 0$.

The goal of the first step, Heuristic Mixing, is to transform an arbitrary permutation into one with the ratio $x : y : 0.5$, where x and y are arbitrary. The method is to apply CX gates a few times until the desired ratio is (nearly) achieved. The following assumption is based on an observation: unstructured or randomly chosen permutations have a ratio close to $x : y : 0.5$.

ASSUMPTION 4.1. *There exists a composite gate consisting of at most four CX and one $C^{n-1}X$ gates that transforms a permutation such that the resulting permutation exhibits the ratio $x : y : 0.5$.*

Here we count $X_iCX_{ij}X_i$, $j \neq i$ (so called negative controlled gates) as CX gate, too. Let us call a composite gate consisting of t CX gates a depth- t composite. For example, counting the meaningful composites, there exist no depth-0 composites, $2n - 2$ depth-1 composites,⁴ at most $(2n - 2) \cdot 4 \binom{n}{2}$ depth-2 composites, and so on. As an algorithm, we may apply depth- t composites one-by-one from $t = 0$ to $t = 4$ until the ratio hits $x : y : 0.5$ exactly, or until all the composites are exhausted. When the latter happens, we choose a permutation among the tried ones that are closest to the desired ratio, and then apply one $C^{n-1}X$ gate to meet $x : y : 0.5$.

We have carried out numerical tests on random samples for $t \in \{0, 1, 2, 3\}$, plotting $\lambda - 2^{n-1}$ (or $|\lambda - 2^{n-1}|$), where λ is the number of interrupting positions in a permutation we can get with depth- t composites that is closest to the desired ratio. As the results in Fig. 5 show, applying a few CX composites likely leads to the desired ratio.

4.2 Preprocessing The output of Heuristic Mixing is a permutation with the ratio $x : y : 0.5$. It is then further processed by an algorithm \mathcal{A}_{pre} to be the ratio $0.5 : 0.5 : 0$.

The idea is simple and we give an example first. Consider the following 4-bit permutation:

$$(3, \mathbf{10}, 14, 6, \mathbf{12}, 2, \mathbf{0}, 15, 5, \mathbf{8}, \mathbf{13}, \mathbf{9}, \mathbf{1}, 4, 7, \mathbf{11}),$$

where eight interrupting positions are highlighted. Move four of them to the left by a sequence of gates $CX_{13}, X_1, CX_{12}, CX_{13}, CX_{241}, CX_{32}$ in order, leading to

$$(\mathbf{13}, \mathbf{9}, \mathbf{1}, \mathbf{10}, 7, 6, 5, \mathbf{8}, \mathbf{0}, 15, 3, 4, 14, \mathbf{11}, \mathbf{12}, 2).$$

Applying another sequence $X_1, X_2, CX_{124}, X_1, X_2$ results in

$$(9, 13, 10, 1, 7, 6, 5, 8, 0, 15, 3, 4, 14, 11, 12, 2).$$

Here we only sketch the mechanism, rather than elaborating on it. Due to the way the interrupting position is defined, it can only be an even number; (hypothetical) change of a ‘single’ row number leads to either $-2, +0$, or $+2$ change to the number of interrupting positions. Since at least two-row numbers are exchanged by the logic gates, the number of interrupting positions can only vary by multiples of 4. Note that an action of $C^{n-1}X$ gate that swaps two-row numbers can change the number of interrupting positions by at most 4. Considering along the line, an action of C^2X gate that swaps 2^{n-2} row numbers can change the number by at most 2^{n-1} . Since the number of interrupting positions of an input permutation is 2^{n-1} , at best a single Toffoli can achieve the desired ratio. However, the best case hardly happens naturally, and we need to manipulate the input before applying the Toffoli

⁴The number of interrupting positions can be varied by CX gate if the gate’s target is the n -th bit. There exist $2n - 2$ such CX gates.

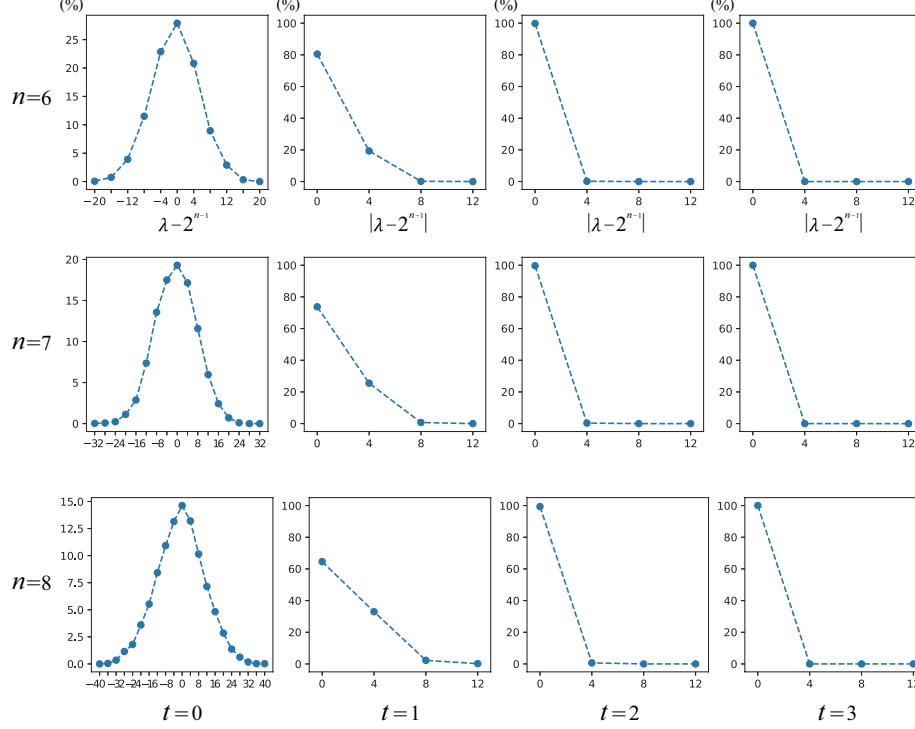


Figure 5: Sampling tests for $n \in \{6, 7, 8\}$ and $t \in \{0, 1, 2, 3\}$. Each figure shows the percentage of samples as a function of $\lambda - 2^{n-1}$ (or $|\lambda - 2^{n-1}|$), where λ is defined in the main text. In other words, figures show how far the permutations are from the ratio $x : y : 0.5$. The margin of error is $\pm 1.55\%$ at 95% confidence level.

gate. It is complicated to give details, but we claim that a procedure similar to repeating CONS and ALLOC one-quarter times of that of $\mathcal{A}'_{\text{red}}$ is enough. In addition to its ability to eliminate interrupting positions, the procedure can also fine-tune the number of normal and inverted positions. In fact, \mathcal{A}_{pre} can turn the ratio $x : y : 0.5$ into $x \pm \Delta_x : y \pm \Delta_y : 0$ where $0 \leq \Delta_x, \Delta_y \leq 0.5$. In the following algorithm \mathcal{A}_{pre} , we may simply set the output ratio to be 0.5:0.5:0.

The procedure is described as follows:

Algorithm 2 \mathcal{A}_{pre}

Input P_n \triangleright The ratio is $x : y : 0.5$ **Output** (P, R) **Procedure**

```
1:  $R \leftarrow ()$ 
2:  $P \leftarrow P_n$ 
3: for  $i$  from 0 to  $2^{n-3} - 1$  do
4:    $\langle a, b \rangle \leftarrow \text{PRE\_PICK}(P, i)$ 
5:    $S_C \leftarrow \text{CONS}(P, i, \langle a, b \rangle)$ 
6:    $(P, R) \leftarrow (P, R) \cdot S_C$ 
7:    $S_A \leftarrow \text{ALLOC}(P, i, a)$ 
8:    $(P, R) \leftarrow (P, R) \cdot S_A$ 
9:  $(P, R) \leftarrow (P, R) \cdot (X_1, X_2, CX_{12n}, X_2, X_1)$ 
10: return  $(P, R)$ 
```

PRE_PICK works in a similar way to PICK, but the output is not a relevant row pair but a certain pair of row numbers at interrupting positions. We will not cover the details on the pair, but with abuse of notation we let $\langle \cdot, \cdot \rangle$ denotes the pair, too. Interested readers may refer to the implemented code [1]. CONS and ALLOC work exactly the same way as in Section 3.4. The number of Toffoli gates involved in Algorithm 2 is bounded by $3 \cdot 2^{n-4} - 1 + \sum_{i=2}^{n-3} (2i - 3) \cdot \binom{n-3}{i}$.

4.3 Generalized Size Reduction Algorithm After the mixing and the preprocessing, the permutation has the ratio $0.5 : 0.5 : 0$. Now the problem can be thought of as two subproblems. We will first construct and allocate only *even blocks* out of row numbers that are in normal positions. Since CONS and ALLOC do not require any controlled gate targeting the n -th bit, the number of normal or inverted positions is conserved by Remark 3.1. Once 2^{n-2} even blocks are left-allocated, the remaining row numbers that are only in inverted positions are constructed and allocated to be *odd blocks*. In the end, we would have 2^{n-2} even blocks on the left half and 2^{n-2} odd blocks on the right half. Applying one CX_{1n} gate completes the size reduction.

The size reduction algorithm is described as follows:

Algorithm 3 \mathcal{A}_{red}

Input P_n ▷ The ratio is 0.5 : 0.5 : 0
Output (P, R) ▷ $P = P_{n-1} \otimes I_2$
Procedure

- 1: $R \leftarrow ()$
- 2: $P \leftarrow P_n$
- 3: **for** i from 0 to $2^{n-2} - 1$ **do** ▷ First Part
- 4: $\langle a, b \rangle \leftarrow \text{N_PICK}(P, i)$ ▷ Pick a normal pair
- 5: $S_C \leftarrow \text{CONS}(P, i, \langle a, b \rangle)$
- 6: $(P, R) \leftarrow (P, R) \cdot S_C$
- 7: $S_A \leftarrow \text{ALLOC}(P, i, a)$
- 8: $(P, R) \leftarrow (P, R) \cdot S_A$
- 9: **for** i from 2^{n-2} to $2^{n-1} - 1$ **do** ▷ Second Part
- 10: $\langle a, b \rangle \leftarrow \text{PICK}(P, i)$
- 11: $S_C \leftarrow \text{CONS}(P, i, \langle a, b \rangle)$
- 12: $(P, R) \leftarrow (P, R) \cdot S_C$
- 13: $S_A \leftarrow \text{ALLOC}(P, i, a)$
- 14: $(P, R) \leftarrow (P, R) \cdot S_A$
- 15: $(P, R) \leftarrow (P, R) \cdot (CX_{1n})$
- 16: **return** (P, R)

The first part deals with row numbers in normal positions to construct and allocate even blocks only. Therefore, N_PICK has to output a relevant row pair in normal positions. Accordingly, the working mechanism of N_PICK is a bit different from that of PICK [1]. Quality bounds given by Proposition 3.1, Lemma 3.1, and Lemma 3.2 are no longer valid in the first part, but overall it only adds 2^{n-3} to the quality bound of $\mathcal{A}'_{\text{red}}$. Details on N_PICK and related quality bound are not discussed here. The second part works exactly the same as Algorithm 1 except now the row numbers are only in inverted positions instead of normal positions.

The additional quality factor, 2^{n-3} , can hardly be met in average instances, and in practice the quality difference between outputs of $\mathcal{A}'_{\text{red}}$ and \mathcal{A}_{red} is negligible.

Note that the point of mixing and preprocessing is to transform an arbitrary permutation so that the result has an appropriate form for the decomposition. If a given permutation is already well-suited for $\mathcal{A}'_{\text{red}}$, or \mathcal{A}_{red} , or something similar, mixing and preprocessing can be skipped.

4.4 Algorithm for Synthesis Combining Heuristic Mixing, Preprocessing, and \mathcal{A}_{red} , one can easily come up with an algorithm for reversible logic circuit synthesis as follows:

Algorithm 4 \mathcal{A}_{syn}

Input P_n **Output** R **Procedure**

```
1:  $R \leftarrow ()$ 
2: for  $i$  from  $n$  to 3 do
3:    $(P_i, R_{\text{mix}}) \leftarrow \mathcal{A}_{\text{mix}}(P_i)$  ▷ Heuristic Mixing
4:    $(P_i, R_{\text{pre}}) \leftarrow \mathcal{A}_{\text{pre}}(P_i)$ 
5:    $(P_i, R_{\text{red}}) \leftarrow \mathcal{A}_{\text{red}}(P_i)$ 
6:    $P_{i-1} \leftarrow \text{REDUCE}(P_i)$  ▷  $P_i = P_{i-1} \otimes I_2$ 
7:    $R \leftarrow R; R_{\text{mix}}; R_{\text{pre}}; R_{\text{red}}$  ▷ Append all
8:  $R \leftarrow R; \text{SEARCH}(P_2)$  ▷ Exhaustive search for  $P_2$ 
9: return  $R$ 
```

Let the quality bound given by Theorem 3.1 be $A(n)$ and the cost Eq. (4.10) be $B(n)$ for an n -bit permutation. The number of Toffoli gates required for the size reduction $P_n \mapsto P_{n-1}$ is upper bounded by

$$A(n) + B(n). \quad (4.11)$$

Therefore the quality bound on the output R of Algorithm 4 is given by $\sum_{x=n}^3 (A(x) + B(x))$.

Time complexity of \mathcal{A}_{mix} can be estimated by counting the number of composites times 2^n , which reads $O(n^7 2^n)$. Preprocessing roughly does one-quarter of what the size reduction performs, thus the time complexity is $O(n 2^{2n-2})$. Time complexities of $\mathcal{A}'_{\text{red}}$ and \mathcal{A}_{red} are the same. In \mathcal{A}_{syn} , the first size reduction step dominates the overall running time because each time the effective size gets smaller, time to transform the permutation becomes exponentially faster. The time complexity of size reduction steps in \mathcal{A}_{syn} is dominated by \mathcal{A}_{red} with $O(n 2^{2n})$, thus the time complexity of \mathcal{A}_{syn} is $O(n 2^{2n})$.

5 Benchmark and Application

Algorithm 4 can be considered as a very basic algorithm that makes use of the size reduction idea. Indeed we have focused on simplifying the algorithm so that the design criteria are as transparently brought out as possible. When it comes to optimizations we have noticed and indeed gone through a number of directions, for example exploiting partial search or statistical properties or undoing and retrying, but then what we confronted at the end was way too complicated descriptions for the algorithm. Most optimization options we have considered earlier thus have been dropped at the end, not only from the paper but also from the source code so that the code can be a natural reflection of the descriptions given in Sections 3 and 4.

Nevertheless, one optimization option is still implemented in our code. Briefly explained, as shown in Subroutines CONS and ALLOC, there is a possibility that a block is already existing at i -th iteration. We call such cases free constructions or free allocations, or simply free blocks without distinction.⁵ To utilize free blocks, let us review how the algorithm works.

In i -th iteration in Algorithm 3, N_PICK or PICK chooses one relevant row pair and it is processed to be a left-allocated block. The resulting permutation may or may not have free blocks in $(i+1)$ -th iteration. Notice at this point that instead of specifying only one relevant row pair for the block, we may try all possible relevant row pairs and rank them based on certain criteria, for example by the

⁵We remark that the chance of having free blocks are not uniform over the iterator i in Algorithm 3.




Figure 6: Quality bounds on size reduction of n -bit permutations (Eq. (4.11)) joined by dashed lines, and the statistical averages of the output qualities denoted by the red bars for 1000 randomly generated instances for each n with $d_n = 0$. Min-Max values are also marked with black bars.

number of free blocks upon the left-allocation of the candidate pair (block). Let us call it depth-0 exhaustive search. Figure 6 shows the number of Toffoli gates resulted from the size reduction of 1000 randomly generated permutations for $n = 4, 5, 6, 7, 8$ employing depth-0 exhaustive search, together with the upper bound for the reference. If computational resources allow, one may try searching for a better pair by considering free blocks upon the construction and allocation of blocks not only at the i -th block-wise position but also up to $(i + d_n(i))$ -th position, where $d_n(i)$ is called a search-depth at i for an n -bit permutation. For example, setting $d_n(i) = 1$ for all $i \in \{0, 1, \dots, 2^{n-1} - 2\}$ is to carry out depth-1 exhaustive search in constructing and allocating each block. The quality of the output circuit gets better as $d_n(i)$ gets larger, but the time complexity becomes worse. Roughly estimated, for a constant $d_n(i) = d$, the time complexity reads $O(n2^{(2+(d+1))n})$.

Since the number of remaining relevant row numbers decreases as the iteration goes, it becomes even easier to do depth- t exhaustive search for larger t in later iterations. In our code, we basically let $d_n(i)$ be a constant and let users control it as an external parameter d_n for each n , except $d_n(i)$ for $i \geq 2^{n-1} - 9$.⁶

Reversible logic circuits for known hard benchmark functions [24, 17] have been synthesized by the algorithm. Comparisons with the previous works are given in Table 1. In ‘best reported’ columns, we have extracted the results with the minimum QC cost listed in the benchmark pages. For d_n options used, see [1].

The algorithm is applied to find reversible circuits for cryptographic S-boxes. Among various S-boxes, one that has attracted the most attention from the research community is undoubtedly AES

⁶It is set such that the last few blocks are exhaustively searched, which has a nontrivial effect on the output quality considering the statistical properties and Lemma 3.1. About the statistical properties, for example, it can be shown that the last four blocks cost the *minimal* number of Toffoli gates if no free block exists upon the construction and allocation of a block in $(2^{n-1} - 5)$ -th iteration. It seems many iteration stages can utilize such statistical properties.

S-box. See, previous works [10, 13, 15, 33] and related references therein. Table 1 in Jaques et al.’s work [13] neatly summarizes a few notable circuit designs from the literature, and we add one more entry to it as in Table 2 with the same Q# options being used as the first entry. The design we have added in the table is first obtained by the algorithm proposed, and then it went through manual optimizations [3, 12] before Q# estimates the resources. A detailed circuit can be found in [1]. Note that we could get an even smaller T-depth at the cost of increasing the number of measurements, but we post the one without any measurements. Here we would point out that our design works *in-place* [8] meaning that it does not need extra bits for writing the outputs. In-place design of the S-box likely leads to a quantum oracle with a minimal number of qubits, but we leave the task of constructing an oracle as one of the future directions to explore. Note however that in Grover’s algorithm it is generally more important to optimize depth than width considering the quadratic speedup [11, 32, 14].

Table 2: Comparison of different circuit designs for AES S-box given in [13] and a circuit obtained from this work. OP and IP in the second column stand for out-of-place and in-place, and the third to the sixth columns read the number of CNOT gates, single qubit Clifford gates, T gates, and measurements. TD and W mean T-depth and the number of logical qubits, respectively.

Source	Type	#CX	#1qC	#T	#M	TD	W
[10]	OP	8683	1023	3854	0	217	44
[15]	OP	818	264	164	41	35	41
[13]	OP	654	184	136	34	6	137
This work	IP	12243	1829	4998	0	578	15

Efficient reversible circuits can be found for S-boxes with algebraic structures, but if an S-box does not have apparent structure, Algorithm 4 can be a good candidate for a solution. We have applied the algorithm to Skipjack [22], KHAZAD [4], and DES S-boxes [20], all of which are known not to have efficient classical implementations other than lookup tables. The results are summarized in Table 3, and the circuits can be found from [1]. To our knowledge, none of them have been analyzed before possibly due to the difficulty of handling unstructured permutation maps. Note that the algorithm is generally applicable to any substitution maps with not too large n , giving rise to in-place logic circuits.

Table 3: Costs of reversible circuits for various S-boxes with no structure. DES has eight different S-boxes, where each takes as input a 6-bit string and outputs a 4-bit string. Two bits of garbage are unavoidable in the reversible implementation of DES S-box.

S-box	#in	#out	#grb	QC	#TOF
Skipjack	8	8	0	5562	791
KHAZAD	8	8	0	5411	794
DES-1	6	4	2	655	87
DES-2	6	4	2	662	91
DES-3	6	4	2	703	94
DES-4	6	4	2	699	98
DES-5	6	4	2	721	97
DES-6	6	4	2	686	97
DES-7	6	4	2	781	107
DES-8	6	4	2	726	99

References

- [1] *Algorithm Implementation*, <https://github.com/ReversibleLogicCircuit/SizeReduction>.
- [2] ARABZADEH, MONA AND SAEEDI, MEHDI, *RCViewer+*, version 2.5, 2013. <http://ceit.aut.ac.ir/QDA/RCV.htm>.
- [3] A. BARENCO, C. H. BENNETT, R. CLEVE, D. P. DIVINCENZO, N. MARGOLUS, P. SHOR, T. SLEATOR, J. A. SMOLIN, AND H. WEINFURTER, *Elementary Gates for Quantum Computation*, Phys. Rev. A, 52 (1995), pp. 3457–3467.
- [4] P. BARRETO AND V. RIJMEN, *The Khazad Legacy-Level Block Cipher*, vol. 97, 2000.
- [5] S. BRAVYI, T. J. YODER, AND D. MASLOV, *Efficient Ancilla-Free Reversible and Quantum Circuits for the Hidden Weighted Bit Function*, arXiv preprint, (2020).
- [6] R. BRYANT, *On the Complexity of VLSI Implementations and Graph Representations of Boolean Functions with Application to Integer Multiplication*, IEEE Transactions on Computers, 40 (1991), pp. 205–213.
- [7] W. J. DALLY AND B. P. TOWLES, *Principles and Practices of Interconnection Networks*, Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2004.
- [8] T. G. DRAPER, S. A. KUTIN, E. M. RAINS, AND K. M. SVORE, *A Logarithmic-Depth Quantum Carry-Lookahead Adder*, Quantum Info. Comput., 6 (2006), p. 351–369.
- [9] O. GOLUBITSKY, S. M. FALCONER, AND D. MASLOV, *Synthesis of the Optimal 4-bit Reversible Circuits*, in Design Automation Conference, 2010, pp. 653–656.
- [10] M. GRASSL, B. LANGENBERG, M. ROETTELER, AND R. STEINWANDT, *Applying Grover’s Algorithm to AES: Quantum Resource Estimates*, in PQCrypto 2016, 2016, pp. 29–43.
- [11] L. K. GROVER, *Quantum Mechanics Helps in Searching for a Needle in a Haystack*, Phys. Rev. Lett., 79 (1997), pp. 325–328.
- [12] Y. HE, M.-X. LUO, E. ZHANG, H.-K. WANG, AND X.-F. WANG, *Decompositions of n -Qubit Toffoli Gates with Linear Circuit Complexity*, International Journal of Theoretical Physics, 56 (2017), pp. 2350–2361.
- [13] S. JAQUES, M. NAEHRIG, M. ROETTELER, AND F. VIRIDIA, *Implementing Grover Oracles for Quantum Key Search on AES and LowMC*, in Advances in Cryptology – EUROCRYPT 2020, A. Canteaut and Y. Ishai, eds., Cham, 2020, Springer International Publishing, pp. 280–310.
- [14] P. KIM, D. HAN, AND K. C. JEONG, *Time-Space Complexity of Quantum Search Algorithms in Symmetric Cryptanalysis: Applying to AES and SHA-2*, Quantum Information Processing, 17 (2018), p. 339.
- [15] B. LANGENBERG, H. PHAM, AND R. STEINWANDT, *Reducing the Cost of Implementing the Advanced Encryption Standard as a Quantum Circuit*, IEEE Transactions on Quantum Engineering, 1 (2020), pp. 1–12.
- [16] Z. LI, H. CHEN, B. XU, X. SONG, AND X. XUE, *An Algorithm for Synthesis of Optimal 3-qubit Reversible Circuits Based on Bit Operation*, in 2008 Second International Conference on Genetic and Evolutionary Computing, 2008, pp. 455–458.
- [17] MASLOV, DMITRI, *Reversible Logic Synthesis Benchmarks Page*. <http://webhome.cs.uvic.ca/%7Edmaslov/>.
- [18] D. M. MILLER, D. MASLOV, AND G. W. DUECK, *A Transformation Based Algorithm for Reversible Logic Synthesis*, in Proceedings of the 40th Annual Design Automation Conference, DAC ’03, New York, NY, USA, 2003, Association for Computing Machinery, p. 318–323.
- [19] M. MÖTTÖNEN, J. J. VARTIAINEN, V. BERGHOLM, AND M. M. SALOMAA, *Quantum Circuits for General Multiqubit Gates*, Phys. Rev. Lett., 93 (2004), p. 130502.
- [20] NATIONAL BUREAU OF STANDARDS, *Data Encryption Standard*, no. 46, 1977.
- [21] M. A. NIELSEN AND I. L. CHUANG, *Quantum Computation and Quantum Information: 10th Anniversary Edition*, Cambridge University Press, 2010.
- [22] NIST, *Skipjack and KEA Algorithm Specifications, Version 2.0*, May 1998.
- [23] NIST, *Advanced Encryption Standard*, FIPS PUB 197, 2001.
- [24] REVLIB, *An Online Resources for Reversible Functions and Circuits*. <http://www.revlib.org/index.php>.
- [25] M. SAEEDI, M. ARABZADEH, M. S. ZAMANI, AND M. SEDIGHI, *Block-Based Quantum-Logic Synthesis*, Quantum Info. Comput., 11 (2011), p. 262–277.
- [26] M. SAEEDI AND I. L. MARKOV, *Synthesis and Optimization of Reversible Circuits— a Survey*, 45 (2013).
- [27] M. SAEEDI, M. S. ZAMANI, M. SEDIGHI, AND Z. SASANIAN, *Reversible Circuit Synthesis Using a Cycle-Based Approach*, J. Emerg. Technol. Comput. Syst., 6 (2010).

- [28] A. SHAFAEI, M. SAEEDI, AND M. PEDRAM, *Reversible Logic Synthesis of k -input, m -output Lookup Tables*, in 2013 Design, Automation Test in Europe Conference Exhibition (DATE), 2013, pp. 1235–1240.
- [29] T. TOFFOLI, *Reversible Computing*, in Automata, Languages and Programming, J. de Bakker and J. van Leeuwen, eds., Berlin, Heidelberg, 1980, Springer Berlin Heidelberg, pp. 632–644.
- [30] G. F. VIAMONTES, I. L. MARKOV, AND J. P. HAYES, *Quantum Circuit Simulation*, Springer Publishing Company, 2009.
- [31] D. V. ZAKABLUKOV, *Application of Permutation Group Theory in Reversible Logic Synthesis*, booktitle="Reversible Computation, Cham, 2016, Springer International Publishing, pp. 223–238.
- [32] C. ZALKA, *Using Grover's Quantum Algorithm for Searching Actual Databases*, Phys. Rev. A, 62 (2000), p. 052305.
- [33] J. ZOU, Z. WEI, S. SUN, X. LIU, AND W. WU, *Quantum Circuit Implementations of AES with Fewer Qubits*, in Advances in Cryptology – ASIACRYPT 2020, S. Moriai and H. Wang, eds., Cham, 2020, Springer International Publishing, pp. 697–726.